

# Computer Science M146, Homework 2

Michael Wu  
UID: 404751542

February 6th, 2018

## Problem 1

a) We have the following points

$x_1$	$x_2$	$y$
-1	-1	-1
-1	1	-1
1	-1	-1
1	1	1

Initializing  $\boldsymbol{\theta} = \langle 0, 0 \rangle$  and  $b = 0$  and running the perceptron algorithm on this list from top to bottom yields  $\boldsymbol{\theta} = \langle 1, 1 \rangle$  and  $b = -1$ . Then

$$y = \begin{cases} 1 & \text{if } \theta_1 x_1 + \theta_2 x_2 + b \geq 0 \\ -1 & \text{if } \theta_1 x_1 + \theta_2 x_2 + b < 0 \end{cases}$$

which separates our training data perfectly. This is not the only unique way to separate our data, as if we chose  $\boldsymbol{\theta} = \langle 2, 1 \rangle$  and  $b = -2$  this would be a valid perceptron as well.

b) We have the following points

$x_1$	$x_2$	$y$
-1	-1	-1
-1	1	1
1	-1	1
1	1	-1

No valid perceptron exists, because this data is not linearly separable.

## Problem 2

a)

$$\begin{aligned}
\frac{\partial J(\boldsymbol{\theta})}{\partial \theta_j} &= -\frac{\partial}{\partial \theta_j} \sum_{n=1}^N \left[ y_n \ln \left( \frac{1}{1 + e^{-\boldsymbol{\theta}^T x_n}} \right) + (1 - y_n) \ln \left( \frac{e^{-\boldsymbol{\theta}^T x_n}}{1 + e^{-\boldsymbol{\theta}^T x_n}} \right) \right] \\
&= -\frac{\partial}{\partial \theta_j} \sum_{n=1}^N \left[ -y_n \ln \left( 1 + e^{-\boldsymbol{\theta}^T x_n} \right) + (1 - y_n) \left( -\boldsymbol{\theta}^T x_n - \ln \left( 1 + e^{-\boldsymbol{\theta}^T x_n} \right) \right) \right] \\
&= -\sum_{n=1}^N \left[ y_n x_{nj} \frac{e^{-\boldsymbol{\theta}^T x_n}}{1 + e^{-\boldsymbol{\theta}^T x_n}} + (1 - y_n) \left( -x_{nj} + x_{nj} \frac{e^{-\boldsymbol{\theta}^T x_n}}{1 + e^{-\boldsymbol{\theta}^T x_n}} \right) \right] \\
&= -\sum_{n=1}^N x_{nj} \left[ y_n \frac{e^{-\boldsymbol{\theta}^T x_n}}{1 + e^{-\boldsymbol{\theta}^T x_n}} + (1 - y_n) \left( \frac{e^{-\boldsymbol{\theta}^T x_n}}{1 + e^{-\boldsymbol{\theta}^T x_n}} - 1 \right) \right] \\
&= -\sum_{n=1}^N x_{nj} \left[ \frac{e^{-\boldsymbol{\theta}^T x_n}}{1 + e^{-\boldsymbol{\theta}^T x_n}} + y_n - 1 \right] \\
&= -\sum_{n=1}^N x_{nj} \left( y_n - \frac{1}{1 + e^{-\boldsymbol{\theta}^T x_n}} \right) \\
&= \sum_{n=1}^N x_{nj} (h_{\boldsymbol{\theta}}(x_n) - y_n)
\end{aligned}$$

## Problem 3

a)

$$\nabla J(\theta_0, \theta_1) = 2 \sum_{n=1}^N [w_n(\theta_0 + \theta_1 x_{n,1} - y_n) \langle 1, x_{n,1} \rangle]$$

b) First we will write out the partial derivative with respect to  $\theta_0$  and set it to zero.

$$\begin{aligned}
\frac{\partial J}{\partial \theta_0} &= 2 \sum_{n=1}^N [w_n(\theta_0 + \theta_1 x_{n,1} - y_n)] = 0 \\
&\sum_{n=1}^N [w_n(\theta_0 + \theta_1 x_{n,1} - y_n)] = 0 \\
\theta_0 \sum_{n=1}^N w_n + \theta_1 \sum_{n=1}^N w_n x_{n,1} - \sum_{n=1}^N w_n y_n &= 0 \\
a\theta_0 + b\theta_1 &= c
\end{aligned}$$

where  $a = \sum_{n=1}^N w_n$ ,  $b = \sum_{n=1}^N w_n x_{n,1}$ , and  $c = \sum_{n=1}^N w_n y_n$ . Next we will write out the partial derivative with respect to  $\theta_1$  and set it to zero.

$$\begin{aligned}
\frac{\partial J}{\partial \theta_1} &= 2 \sum_{n=1}^N [x_{n,1} w_n (\theta_0 + \theta_1 x_{n,1} - y_n)] = 0 \\
&\sum_{n=1}^N [x_{n,1} w_n (\theta_0 + \theta_1 x_{n,1} - y_n)] = 0 \\
\theta_0 \sum_{n=1}^N x_{n,1} w_n + \theta_1 \sum_{n=1}^N w_n x_{n,1}^2 - \sum_{n=1}^N x_{n,1} w_n y_n &= 0 \\
b\theta_0 + d\theta_1 &= e
\end{aligned}$$

where  $d = \sum_{n=1}^N w_n x_{n,1}^2$  and  $e = \sum_{n=1}^N x_{n,1} w_n y_n$ . Then solving this system of equations gives

$$\begin{aligned}
\theta_0 &= \frac{e - d\theta_1}{b} \\
\frac{a}{b}(e - d\theta_1) + b\theta_1 &= c \\
\frac{ae}{b} + \left(b - \frac{ad}{b}\right)\theta_1 &= c \\
\frac{b^2 - ad}{b}\theta_1 &= c - \frac{ae}{b} \\
\theta_1 &= \frac{bc - ae}{b^2 - ad}
\end{aligned}$$

$$\begin{aligned}\theta_0 &= \frac{e - d \frac{bc-ae}{b^2-ad}}{b} \\ \theta_0 &= \frac{eb^2 - ade - bcd + ade}{b^3 - abd} \\ \theta_0 &= \frac{eb - cd}{b^2 - ad}\end{aligned}$$

Then we have a minimum at

$$(\theta_0, \theta_1) = \frac{1}{b^2 - ad} \langle eb - cd, bc - ae \rangle$$

## Problem 4

a) Assume that our data set is finite, so we have for all  $(x_i, y_i) \in D$  there exists a  $w$  and  $\theta$  such that

$$y_i(w^T x_i + \theta) > l$$

where  $l$  is the margin between our separating line and the closest point in our data set. Then we can multiply our weight vector  $w$  and our threshold value  $\theta$  by  $\frac{1}{l}$  to get

$$y_i \left( \frac{1}{l} w^T x_i + \frac{\theta}{l} \right) > 1$$

Thus we can find a weight vector and threshold value that creates an optimal solution to the linear program with  $\delta = 0$ .

b) If we have an optimal solution to the linear problem, we have

$$y_i (w^T x_i + \theta) > 1$$

for all  $(x_i, y_i) \in D$ . In order for this to happen,  $y_i$  must always have the same sign as  $w^T x_i + \theta$ . Thus

$$y_i = \begin{cases} 1 & \text{if } w^T x_i + \theta \geq 0 \\ -1 & \text{if } w^T x_i + \theta < 0 \end{cases}$$

for all  $(x_i, y_i) \in D$ , and so  $D$  is linearly separable.

c) If  $0 < \delta < 1$ , then our data set is linearly separable, using the same argument as above. Otherwise we cannot say anything about if our data set is linearly separable.

d) An optimal solution is  $w = 0$  and  $\theta = 0$ . Then

$$y_i(w^T x_i + \theta) \geq 0$$

for any data set. The problem with this is that our resulting linear separator tells us nothing about our data set, as everything is zero.

e) A possible optimal solution is  $w = \langle 1, 1, 1 \rangle$  and  $\theta = 0$ .

## Problem 5

a) The data seems to be negatively correlated, with most points in the bottom right and top left. I believe linear regression will do a relatively good job of fitting the data.




Figure 1: Plot of training data.

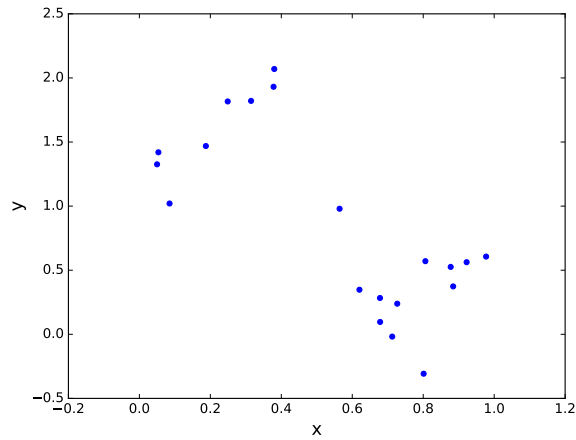


Figure 2: Plot of testing data.

b) I implemented this using the following code

```
m = self.m_
Phi = np.zeros((n,m+1))
for i in range(0,n):
    val=[1]
    index=[(m+1)*i]
    for j in range(0,m):
        val.append(val[j]*X.flat[i])
        index.append(index[j]+1)
    np.put(Phi, index, val)

return Phi
```

This took care of creating feature vectors for linear regression as well as polynomial regression.

c) I simply computed the dot product of each feature vector with the weights in order to make a prediction for  $y$ .

d) The results of the different  $\eta$  values are as follows.

$\eta$	$\theta_0$	$\theta_1$	Iterations	$J(\theta)$	Runtime
0.0001	2.2704	-2.4606	10000	4.0863	1.060399
0.001	2.4464	-2.8163	7021	3.9125	0.794983
0.01	2.4464	-2.8163	765	3.9125	0.079679
0.0407	$-9.40 * 10^{18}$	$-4.65 * 10^{18}$	10000	$2.71 * 10^{39}$	1.042172
$\frac{1}{1+k}$	2.4464	-2.8163	1357	3.9125	0.142043
Closed Form	2.4464	-2.8163	N/A	3.9125	0.000560

The larger step sizes converged quicker, until  $\eta = 0.0407$  when the step size became too large.

e) The closed form solution is

$$\theta = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

The coefficients obtained match the ones obtained with gradient descent, though the runtime was about two to four orders of magnitude faster.

f) It took 1357 iterations and 0.142043 seconds for our algorithm with our proposed learning rate to converge.

g) I already implemented this in part b.

h) We use RMSE because  $J(\theta)$  increases with more training data, which we do not want because it makes it harder to judge how accurate our fit is. RMSE is an average. Additionally,  $J(\theta)$  scales with the square of our error, so we want to take the square root in order to make our error scale more linearly.

i) The polynomial of degree  $m = 5$  best fits our data, with the smallest test RMSE of 0.355137742884. There is evidence of underfitting on the left of the plot, because both training and testing errors are higher there. There is evidence of overfitting on the right, as our training error decreases and our test error increases there.




Figure 3: Plot of how well polynomial regressions of varying degrees fits our data.