

Computer Science M146, Homework 3

Michael Wu
UID: 404751542

February 27th, 2018

Problem 1

The VC dimension of H is 3. An example of 3 points x such that $x \in R$ that can be shattered are $x_1 = -1$, $x_2 = 0$, and $x_3 = 1$. Then the following table shows how we can shatter these points, with $h \in H$ being the classifier used.

x_1 label	x_2 label	x_3 label	h
0	0	0	$\text{sgn}(-1)$
0	0	1	$\text{sgn}(x - 0.5)$
0	1	0	$\text{sgn}(-x^2 + 0.5)$
0	1	1	$\text{sgn}(x + 0.5)$
1	0	0	$\text{sgn}(-x - 0.5)$
1	0	1	$\text{sgn}(x^2 - 0.5)$
1	1	0	$\text{sgn}(-x + 0.5)$
1	1	1	$\text{sgn}(1)$

For any set $S = \{x_1, x_2, x_3, x_4\}$ of four points where $x_1 < x_2 < x_3 < x_4$, we cannot shatter S if we label $x_1 = 1$, $x_2 = 0$, $x_3 = 1$, $x_4 = 0$. This is because any classifier in H at most splits R into three distinct regions where the quadratic $ax^2 + bx + c$ is above or below zero, and we have 4 distinct regions of classification in our set. So our hypothesis space cannot shatter S because no h can correctly classify this training set.

Problem 2

$$\begin{aligned}
K_\beta(\mathbf{x}, \mathbf{z}) &= (1 + \beta \mathbf{x} \cdot \mathbf{z})^3 \\
&= (1 + \beta(x_1 z_1 + \dots + x_D z_D))^3 \\
&= 1 + 3\beta \sum_{i=1}^D x_i z_i + 3\beta^2 \left(\sum_{i=1}^D x_i z_i \right)^2 + \beta^3 \left(\sum_{i=1}^D x_i z_i \right)^3 \\
&= 1 + 3\beta \sum_{i=1}^D x_i z_i + 3\beta^2 \sum_{i,j=1}^D x_i z_i x_j z_j + \beta^3 \sum_{i,j,k=1}^D x_i z_i x_j z_j x_k z_k \\
&= \left\langle 1, \sqrt{3}\beta x_i \Big|_{i=1}^D, \sqrt{3}\beta x_i x_j \Big|_{i,j=1}^D, \beta^{\frac{3}{2}} x_i x_j x_k \Big|_{i,j,k=1}^D \right\rangle \\
&\quad \cdot \left\langle 1, \sqrt{3}\beta z_i \Big|_{i=1}^D, \sqrt{3}\beta z_i z_j \Big|_{i,j=1}^D, \beta^{\frac{3}{2}} z_i z_j z_k \Big|_{i,j,k=1}^D \right\rangle
\end{aligned}$$

Thus we have

$$\phi_\beta(\mathbf{x}) = \left\langle 1, \sqrt{3}\beta x_i \Big|_{i=1}^D, \sqrt{3}\beta x_i x_j \Big|_{i,j=1}^D, \beta^{\frac{3}{2}} x_i x_j x_k \Big|_{i,j,k=1}^D \right\rangle$$

The kernel $K(\mathbf{x}, \mathbf{z}) = (1 + \mathbf{x} \cdot \mathbf{z})^3$ is equivalent to setting $\beta = 1$, and corresponds to the feature map

$$\phi_1(\mathbf{x}) = \left\langle 1, \sqrt{3}x_i \Big|_{i=1}^D, \sqrt{3}x_i x_j \Big|_{i,j=1}^D, x_i x_j x_k \Big|_{i,j,k=1}^D \right\rangle$$

The parameter β scales the features up and down by a constant. It effectively replaces the vector \mathbf{x} by $\sqrt{\beta}\mathbf{x}$ such that $\phi_\beta(\mathbf{x}) = \phi_1(\sqrt{\beta}\mathbf{x})$.

Problem 3

a) We wish to find $\mathbf{w}^* = \langle w_1, w_2 \rangle$ such that we minimize $\frac{1}{2}\sqrt{w_1^2 + w_2^2}$ subject to $w_1 + w_2 \geq 1$ and $-w_1 \geq 1$. Equivalently we would like to find the point closest to the origin such that $w_2 \geq 1 - w_1$ and $w_1 \leq -1$. Thus we get

$$\mathbf{w}^* = \langle -1, 2 \rangle$$

and the margin is $\gamma = \frac{1}{\sqrt{5}}$.

b) We wish to find $\mathbf{w}^* = \langle w_1, w_2 \rangle$ such that we minimize $\frac{1}{2}\sqrt{w_1^2 + w_2^2}$ subject to $w_1 + w_2 + b \geq 1$ and $-w_1 - b \geq 1$. The classifier changes by making the decision boundary a horizontal line that crosses $(0, \frac{1}{2})$, and the margin will increase compared to our previous results. We get

$$\mathbf{w}^* = \langle 0, 2 \rangle$$

and

$$b^* = -1$$

Our margin is $\gamma = \frac{1}{2}$. This makes sense because our two points $\mathbf{x}_1 = (1, 1)$ and $\mathbf{x}_2 = (1, 0)$ have a distance of 1 between them, so our margin is exactly half of this distance. Our solution with offset has a higher margin γ and smaller magnitude of \mathbf{w}^* than our solution without offset.

Problem 4

4.1

a) I created the dictionary using the following code

```
word_list = {}
count = 0
with open(infile, 'rU') as fid :
    for line in fid:
        wordListLine = extract_words(line)
        for word in wordListLine:
            if word not in word_list:
                word_list[word]=count
                count+=1
return word_list
```

b) I extracted the feature vectors using the following code

```
with open(infile, 'rU') as fid :
    lineNum=0
    for line in fid:
        wordListLine = extract_words(line)
        for word in wordListLine:
```

```

        feature_matrix[lineNum,word_list[word]]=1
    lineNum+=1
return feature_matrix

```

c) I split the features and labels into train and test sets using the following code

```

trainX = X[0:560]
trainy = y[0:560]
testX = X[560:630]
testy = y[560:630]

```

d) The feature matrix has the dimensions (630,1811). The trainX set has dimensions (560,1811), the trainy set has dimensions (560,1), the testX set has dimensions (70,1811), and the testy set has the dimensions (70,1).

4.2

a) I implemented performance using the following code

```

score = 0
if metric=="accuracy":
    score=metrics.accuracy_score(y_true,y_label)
if metric=="f1-score":
    score=metrics.f1_score(y_true,y_label)
if metric=="auROC":
    score=metrics.roc_auc_score(y_true,y_pred)
return score

```

b) I implemented cv_performance using the following code

```

return cross_val_score(clf, X, y, scoring=metric, cv=kf).mean()

```

and in main I added

```

kf=StratifiedKFold(trainy, 5)

```

It is beneficial to use a stratified K-fold so that the percentage of positive and negative reviews are the same across folds because this ensures that our

cross validation gets a good representation of our training data in each fold. Otherwise we may accidentally divide the folds such that one gets a small number of negative or positive reviews, making our classifier performance inaccurate.

c) I implemented `select_param_linear` with the following code

```
best=0
cBest=0
for c in C_range:
    score=cv_performance(SVC(kernel="linear",C=c),X,y,kf,metric)
    print "C="+str(c)+" score="+str(score)
    if score>best:
        best=score
        cBest=c
return cBest
```

d) My results were as follows

C	Accuracy	F1-score	AUROC
10^{-3}	0.7089	0.8297	0.8105
10^{-2}	0.7107	0.8306	0.8111
10^{-1}	0.8060	0.8755	0.8576
10^0	0.8146	0.8749	0.8712
10^1	0.8182	0.8766	0.8696
10^2	0.8182	0.8766	0.8696
best C	10^1	10^1	10^0

4.3

a) I chose the hyperparameter $c = 10$ and trained my classifier using the following code

```
clf=SVC(kernel="linear", C=10)
clf.fit(trainX, trainy)
```

b) I implemented `performance_test` using the following code

```
y_pred=clf.decision_function(X)
print "Performance test with metric "+str(metric)+":"
return performance(y, y_pred, metric)
```

c) My results were as follows

Metric	Score
Accuracy	0.7429
F1-score	0.4375
AUROC	0.7454